

# Supplementary Information for RankCloak: Hiding Synthetic Cryptographic Artifacts in Plain English with Language Models

Alexander V. Mantzaris<sup>1,\*</sup>

<sup>1</sup>School of Data, Mathematical, and Statistical Sciences, University of Central Florida, Orlando, FL, United States of America

\*Correspondence: alexander.mantzaris@ucf.edu

## ABSTRACT

This Supplementary Information documents the extended RankCloak methods and the detailed evidence supporting the revised manuscript. It specifies deterministic rank ordering, payload codecs, segmented forced spans, replay contracts, filtering, corpus construction, prompt and model identities, and work-unit accounting. It reports recovery strata, multilingual secondary tests, controlled transmission perturbations, automated readability and model-based quality diagnostics, reviewer-relevant ablation contrasts and all topic contrasts, neural detector splits and leakage sensitivity, capacity–quality relationships, computational overhead, and worked examples. Evidence classifications and unresolved boundaries are retained explicitly. The material does not constitute human evaluation, deployment validation, broad multilingual naturalness evidence, cryptographic-security proof, or a claim of robust visible-text transmission.

## Supplementary Note S1: Extended RankCloak methods

The main Methods contains the rank definition, Calgacus transformation, complete payload-codec equations, and the full non-segmented, segmented, and natural-tail algorithms. This note records implementation details and edge cases that extend those central presentations without duplicating their equations or pseudocode.

### 5 Direct-subword framing and bounded-codec metadata

The direct-subword arm tokenizes the original UTF-8 payload bytes with special-token insertion disabled and computes its ranks under the model’s empty/BOS payload context. A preflight detokenizes the token sequence and requires the original payload bytes to be an exact suffix. The only accepted extra bytes are an explicitly retained leading-space prefix; the record stores its byte length, base64 value, and SHA-256 digest. A non-space prefix, any suffix transformation, or any payload-byte mismatch makes the representation unavailable before cover generation. During inverse transcoding, recovered ranks regenerate the payload-side token IDs, the saved prefix is removed only after exact prefix verification, and the recovered bytes and digest must equal the original. This framing correction prevents tokenization-level equality from being mistaken for recovery of the declared artifact.

For the ASCII  $B = 8$  and  $B = 16$  conditions, metadata stores the alphabet size, bits per symbol, exact source-byte length, and the number of zero bits appended at the end of the complete serialized bitstream. The decoder validates every recovered rank, reconstructs the padded bitstream, removes only the recorded terminal padding, and returns exactly the declared number of bytes. The  $B = 8$  codec therefore uses  $\lceil 8m/3 \rceil$  ranks for  $m$  serialized bytes rather than independently padding each byte to three radix-8 digits. The hexadecimal condition records the exact canonical character length and rejects ranks outside 1–16. It is never applied to base64 or hyphenated UUID text.

### 20 Candidate masks and deterministic ordering

The implementation converts logits to 64-bit numerical scores and orders eligible candidates with a lexicographic sort whose primary key is decreasing score and secondary key is increasing token ID. Rank lookup counts candidates with a strictly greater score plus equal-score candidates having a lower ID. This realizes the tie contract in the main Methods without a stochastic operation.

25 The canonical `safe_text_filter_v1` is constructed once per model and excludes empty or replacement-character pieces, disallowed controls, and declared code, markup, URL, escape, and formatting fragments. The no-filter condition uses the complete candidate vocabulary. The stricter `roundtrip_stable_filter_v1` additionally requires that each isolated token piece detokenize and retokenize to that same single ID. This isolated condition does not imply contextual visible-text

stability. Its Mistral mask was empty, so the associated cells and their dependent analyses are recorded as unavailable rather than failed. Mask identity, allowed-token count, and digest are retained with each trial.

### Segment records, lead-ins, and tail policies

Each segment record stores its rank slice, prompt identity and tokenized context, exact lead-in, forced, tail, and full token-ID arrays, token-role mask, and half-open forced offsets. Single-topic segmentation reuses the assigned prompt; multi-topic segmentation rotates deterministically through the six prompt categories while preserving the assigned template index. Neither segment order nor boundary metadata is authenticated by RankCloak itself.

A lead-in is a serial rank-1 prefix generated before a forced span. Saved-token replay supplies those exact IDs; the separate greedy-lead-in mode regenerates them and checks equality before decoding. Because a lead-in changes every forced-token history, it is not interchangeable with a tail. Lead-in lengths 2, 4, 8, 16, and 32 were compared with the no-lead-in condition.

The fixed 20–60-token sentence-tail policy is Algorithm 3 in the main Methods. The completed study also froze `dynamic_completion_v1`: after at least eight greedy tokens, stopping requires terminal punctuation or a double newline, balanced parentheses, brackets, braces, and quotation marks, and no declared trailing connective or punctuation fragment. Reaching 256 tokens without satisfying the rule is an explicitly censored emergency stop. The ablation compares dynamic stopping with the fixed sentence-tail policy and with no tail. Every tail token retains the role `tail` and is excluded from rank reconstruction.

### Replay modes and structured endpoints

The main Algorithms 1 and 2 define saved-token exact-copy recovery. Detokenized-text replay instead renders the complete visible message, invokes the embedded tokenizer again after stripping one automatically inserted BOS token when applicable, and applies the saved forced offsets. This boundary rule is a diagnostic; it does not search for alternate tokenizations or infer a payload span from prose. Transformed-text replay applies one frozen perturbation before this step. Cross-model replay uses a different declared model/tokenizer pair.

For every mode, the record distinguishes exact forced-rank replay, exact representation recovery, and exact serialized-payload recovery. The primary endpoint requires the last of these, including byte length and SHA-256 identity. Failures retain the first segment and forced position where available, expected and observed token IDs and ranks, the context digest, and boundary offsets. Payload-group and model–payload-group endpoints require all constituent trials to pass. Natural tails never enter any recovery endpoint.

## Supplementary Note S2: Corpus, models, prompts, and work-unit accounting

### Deterministic cryptographic corpus

All 480 artifacts are public deterministic test vectors generated from declared seeds and procedures. They exercise standard formats without representing operational secrets, user credentials, or deployed traffic.

| Payload class       | Count | Algorithm-defined serialization                                         |
|---------------------|-------|-------------------------------------------------------------------------|
| SHA-256 digest      | 60    | Lowercase hexadecimal 256-bit digest                                    |
| HMAC-SHA256 tag     | 60    | Lowercase hexadecimal keyed-hash output under deterministic test inputs |
| Deterministic nonce | 60    | 96-bit nonce serialization                                              |
| 128-bit token       | 60    | Deterministically generated lowercase hexadecimal token                 |
| UUIDv4              | 60    | Standards-conforming version/variant bits and canonical UUID display    |
| AES-256-GCM         | 60    | Base64 serialization of deterministic authenticated-encryption output   |
| ChaCha20–Poly1305   | 60    | Base64 serialization of deterministic AEAD output                       |
| Ed25519 signature   | 60    | Base64 serialization of a deterministic signature test output           |

**Table S1.** Cryptographic test-vector corpus. Counts are balanced by class. “Real” refers to outputs of the named algorithms, not to private operational data.

### Pinned model and tokenizer identities

All models used the `llama_cpp` backend, one model loaded at a time, embedded GGUF tokenizers, rank-replay batch and microbatch sizes of one, and Q4\_K\_M quantization. The retained manifests require exact artifact hashes and backend/library identities.

| Model ID                        | GGUF file                            | Size (bytes)  | SHA-256                                                          |
|---------------------------------|--------------------------------------|---------------|------------------------------------------------------------------|
| llama3_8b_instruct_q4_k_m       | Meta-Llama-3-8B-Instruct.Q4_K_M.gguf | 4,920,734,272 | 86c8ea6c8b755687d0b723176fcd0b2411ef80533d23e2a5030f845d13ab2db7 |
| qwen2.5_7b_instruct_q4_k_m      | Qwen2.5-7B-Instruct-Q4_K_M.gguf      | 4,683,074,240 | 65b8fcd92af6b4fefa935c625d1ac27ea29dcb6ee14589c55a8f115ceaaa1423 |
| mistral_7b_instruct_v0_3_q4_k_m | Mistral-7B-Instruct-v0.3-Q4_K_M.gguf | 4,372,812,000 | 1270d22c0fbb3d092fb725d4d96c457b7b687a5f5a715abe1e818da303e562b6 |

**Table S2.** Executed model artifacts. Repository revisions for Qwen and Mistral and all upstream identities are retained in the model manifest. The three configured GGUF files are not present in the current checkout, so no fresh local generation replay was attempted for this revision.

### English prompt identities

The 18 templates were fixed before the completed primary matrix. Multi-topic segmentation rotated through this schedule; single-topic segmentation reused its declared topic.

**Table S3.** Complete English prompt schedule.

| Category                   | Prompt ID                   | Exact prompt text                                                                                                                              |
|----------------------------|-----------------------------|------------------------------------------------------------------------------------------------------------------------------------------------|
| Casual conversation        | casual_weekend_chat         | Write a friendly message to a friend about making relaxed plans for the weekend. Use ordinary conversational English and complete thoughts.    |
| Casual conversation        | casual_hobby_chat           | Write a casual message to a friend describing a hobby you have been enjoying lately. Keep the tone natural, warm, and specific.                |
| Casual conversation        | casual_neighborhood_chat    | Write an informal message to a neighbor about something pleasant noticed nearby. Use clear everyday language and finish the message naturally. |
| Professional communication | professional_project_update | Write a concise professional project update for colleagues, explaining recent progress, the next step, and one practical consideration.        |
| Professional communication | professional_schedule_note  | Write a polite workplace message proposing a meeting time and briefly explaining what the discussion should accomplish.                        |
| Professional communication | professional_process_email  | Write a short professional email clarifying a routine process and inviting colleagues to raise questions. Use complete, neutral prose.         |
| Forum Q&A                  | forum_practical_answer      | Write a helpful answer to an online forum question about solving a common household problem. Explain the reasoning in plain English.           |
| Forum Q&A                  | forum_purchase_advice       | Write a balanced forum reply to someone comparing two everyday purchases. Mention useful tradeoffs without sounding promotional.               |
| Forum Q&A                  | forum_learning_answer       | Write a supportive question-and-answer style response to a beginner learning a practical skill. Give concrete, cautious advice.                |
| Recipe/procedure           | procedure_simple_meal       | Write clear prose instructions for preparing a simple weeknight meal. Explain the order of the steps and how to tell when it is ready.         |
| Recipe/procedure           | procedure_home_task         | Write an easy-to-follow explanation of a routine home maintenance task. Include preparation, the main steps, and a sensible final check.       |
| Recipe/procedure           | procedure_garden_task       | Write practical instructions for a basic gardening task. Use natural sentences, explain why the steps matter, and end cleanly.                 |
| Personal narrative/blog    | narrative_small_discovery   | Write a short personal blog passage about discovering something useful during an ordinary day. Keep the reflection concrete and understated.   |
| Personal narrative/blog    | narrative_local_walk        | Write a first-person narrative about a walk through a familiar place and one detail that changed the writer's perspective.                     |
| Personal narrative/blog    | narrative_learning_moment   | Write a brief personal story about learning from a minor mistake. Use coherent, natural prose and a restrained conclusion.                     |
| Factual/explanatory        | explain_weather_pattern     | Explain a familiar weather pattern to a general audience. Connect cause and effect clearly without using equations or specialist jargon.       |

| Category            | Prompt ID       | Exact prompt text                                                    |
|---------------------|-----------------|----------------------------------------------------------------------|
| Factual/explanatory | explain_        | Explain how an everyday public system works for a curious reader.    |
|                     | everyday_system | Use accurate, neutral language and a logically complete explanation. |
| Factual/explanatory | explain_        | Explain a basic biological process in accessible prose. Define the   |
|                     | biological_     | central idea, describe the sequence, and avoid unsupported claims.   |
|                     | process         |                                                                      |

## Experimental matrices and reconciliation

| Matrix             | Declared units | Successful | Unavailable | Scope                                                                     |
|--------------------|----------------|------------|-------------|---------------------------------------------------------------------------|
| Primary            | 14,400         | 14,400     | 0           | Six RankCloak protocol/codec conditions plus ordinary controls            |
| Ablation           | 1,872          | 1,824      | 48          | Lead-in, segment size, filter, and tail policy                            |
| Robustness         | 3,744          | 3,408      | 336         | Replay, raw transformations, limited mitigation, and cross-model mismatch |
| Multilingual       | 1,152          | 1,152      | 0           | Spanish and Simplified Chinese payload/control generation                 |
| Held-out evaluator | 17,280         | 17,232     | 48          | Source and held-out model scoring                                         |
| Detector           | 56             | 56         | 0           | Two architectures across 28 frozen splits                                 |
| Total              | 38,504         | 38,072     | 432         | All declared work reconciled; zero execution failure; zero remaining      |

**Table S4.** Complete work-unit accounting. “Unavailable” means the declared source condition or compatible retained output did not exist; it is excluded from estimands and is not an observed experimental failure.

The 432 unavailable units comprise 48 Mistral roundtrip-stable-filter ablation units, 336 Mistral limited-canonicalization robustness units (seven conditions times 48), and 48 corresponding held-out-evaluator units. The root condition was an empty isolated roundtrip vocabulary. All other matrices completed. Primary work comprised 1,440 trials each for direct subword, ASCII  $B = 8$ , and ASCII  $B = 16$ , and 720 each for non-segmented hex, segmented single-topic hex, and segmented multi-topic hex. The detector corpus and fits are accounted as declared detector work, not as additional generation.

## Supplementary Note S3: Detailed recovery and multilingual results

| Stratum                            | Success | Total | Interpretation                                |
|------------------------------------|---------|-------|-----------------------------------------------|
| Llama 3 8B, all payload protocols  | 2,160   | 2,160 | Saved-token exact serialized-payload recovery |
| Mistral 7B, all payload protocols  | 2,160   | 2,160 | Same endpoint                                 |
| Qwen 2.5 7B, all payload protocols | 2,160   | 2,160 | Same endpoint                                 |
| Direct subword                     | 1,440   | 1,440 | All three models                              |
| ASCII byte, $B = 8$                | 1,440   | 1,440 | All three models; rank ceiling 8              |
| ASCII byte, $B = 16$               | 1,440   | 1,440 | All three models; rank ceiling 16             |
| Hex nibble, non-segmented          | 720     | 720   | Eligible hex-like payloads                    |
| Hex nibble, segmented single-topic | 720     | 720   | Forced spans decoded; tails ignored           |
| Hex nibble, segmented multi-topic  | 720     | 720   | Forced spans decoded; tails ignored           |

**Table S5.** Primary recovery by model and protocol. Every row uses the same exact serialized-payload endpoint. These overlapping summaries are descriptive strata, not additional independent trials.

At the frozen trial endpoint, 6,480/6,480 trials recovered, with Wilson 95% CI [0.9994075335, 1.0000000000]. At the strict payload endpoint, a payload passed only when every observed model, prompt, and protocol trial passed: 480/480, CI [0.9920605009, 1.0000000000]. At the model–payload endpoint, the result was 1,440/1,440, CI [0.9973394178, 1.0000000000]. Exact rank replay was retained as a diagnostic; the confirmatory outcome was the original serialized payload bytes and digest.

| Language           | Trial success | Payload-group success | Trial Wilson 95% CI        | Scope                                                |
|--------------------|---------------|-----------------------|----------------------------|------------------------------------------------------|
| Spanish            | 288/288       | 48/48                 | 0.986837–1.000000          | Secondary saved-token exact-copy recovery            |
| Simplified Chinese | 288/288       | 48/48                 | 0.986837–1.000000          | Secondary saved-token exact-copy recovery            |
| Combined           | 576/576       | 96/96                 | Descriptive combined total | No human naturalness or broad multilingual inference |


**Table S6.** Secondary multilingual recovery. Each language’s payload-group Wilson interval is 0.925900–1.000000.

These results do not test translation, code switching, cross-language decoding, or native-speaker judgments. They also remain conditional on the same model/tokenizer artifacts and exact saved tokens. The GGUF artifacts are absent from the current checkout; consequently the retained results can be integrity-checked against their manifests, but fresh generation or rank replay cannot be reproduced locally from this checkout alone.

80

## Supplementary Note S4: Robustness and failure analysis

### Exact-copy fragility



**Figure S1. Complete recovery matrix and failure diagnostics.** The upper panels distinguish saved-token, visible-text, transformed-text, limited canonicalization, and cross-model replay. Intervals use source-cover Wilson calculations rather than treating two cross-model decoders as independent covers. The lower panels summarize first-divergence categories. Those categories identify the first recorded mismatch and are descriptive, not causal proof. Unavailable Mistral limited-mitigation cells are excluded rather than counted as failures.

**Table S7.** Complete principal robustness recovery table.

| Family/replay                       | Condition  | Success | Total | Wilson 95% CI     |
|-------------------------------------|------------|---------|-------|-------------------|
| Replay: saved IDs                   | Unmodified | 144     | 144   | 0.974016–1.000000 |
| Replay: detokenize/retokenize       | Unmodified | 88      | 144   | 0.529589–0.686859 |
| Replay: greedy lead-in regeneration | Unmodified | 144     | 144   | 0.974016–1.000000 |

| Family/replay            | Condition                               | Success | Total                   | Wilson 95% CI                        |
|--------------------------|-----------------------------------------|---------|-------------------------|--------------------------------------|
| Raw transmission         | Unmodified                              | 144     | 144                     | 0.974016–1.000000                    |
| Raw transmission         | Final 10% tail-only truncation          | 144     | 144                     | 0.974016–1.000000                    |
| Raw transmission         | Unicode NFKC                            | 82      | 144                     | 0.487804–0.647476                    |
| Raw transmission         | Quote conversion                        | 56      | 144                     | 0.313141–0.470411                    |
| Raw transmission         | Token deletion                          | 34      | 144                     | 0.174168–0.311768                    |
| Raw transmission         | Whitespace trim                         | 28      | 144                     | 0.138095–0.266672                    |
| Raw transmission         | Whitespace collapse                     | 21      | 144                     | 0.097405–0.212667                    |
| Raw transmission         | Character deletion                      | 18      | 144                     | 0.080551–0.188937                    |
| Raw transmission         | Character insertion                     | 18      | 144                     | 0.080551–0.188937                    |
| Raw transmission         | Character substitution                  | 16      | 144                     | 0.069559–0.172872                    |
| Raw transmission         | LF-to-CRLF line endings                 | 16      | 144                     | 0.069559–0.172872                    |
| Raw transmission         | Markdown blockquote copy/paste          | 0       | 144                     | 0.000000–0.025984                    |
| Raw transmission         | Deterministic paraphrase                | 0       | 144                     | 0.000000–0.025984                    |
| Cross-model retokenized  | Unmodified, two alternate decoders      | 0       | 288 outcomes/144 covers | source-cover CI<br>0.000000–0.025984 |
| Limited canonicalization | Unmodified                              | 96      | 96                      | 0.961524–1.000000                    |
| Limited canonicalization | Each of six available transformed cells | 0       | 96 per cell             | 0.000000–0.038476                    |

The limited canonicalization cells covered line endings, markdown copy/paste, quote conversion, NFKC, whitespace collapse, and whitespace trim. Each had 96 observations and 48 unavailable Mistral source conditions. The all-zero transformed results do not support a mitigation claim.

| Requested class             | Coverage | Frozen proxy                       | Boundary                                            |
|-----------------------------|----------|------------------------------------|-----------------------------------------------------|
| Whitespace changes          | Direct   | Trim, collapse, line endings       | Separate conditions; not a pooled estimand          |
| Leading/trailing whitespace | Direct   | Trim                               | Addition was not tested                             |
| Unicode normalization       | Direct   | NFKC                               | Other forms untested                                |
| Quotation marks             | Direct   | Straight-to-smart conversion       | Reverse/locale-specific forms untested              |
| Copy/paste                  | Direct   | Fixed markdown blockquote pipeline | Not an application/OS survey                        |
| Line wrapping               | Partial  | LF-to-CRLF proxy                   | No visual reflow/insertion test                     |
| Markdown/email quoting      | Partial  | Markdown blockquote                | No client headers, nesting, or reflow               |
| Sentence boundaries         | Partial  | Broad paraphrase                   | Boundary edits not isolated                         |
| Punctuation changes         | Untested | None                               | Alphanumeric substitution is not a punctuation test |
| Arbitrary prefix/suffix     | Untested | None                               | Blockquote markup is not general insertion          |
| Case conversion             | Untested | None                               | No estimate available                               |


**Table S8.** Perturbation-coverage inventory: five directly tested, three partially represented, and three untested requested classes.

85 Failure summaries classify detokenization/retokenization divergence, whitespace/markup tokenization divergence, quotation divergence, Unicode normalization divergence, content-edit divergence, semantic rewrite divergence, and model/tokenizer identity mismatch. For example, a mismatch can appear at the first forced position because a transformed prefix changes the context, even when later visible words are unchanged. The recorded first-difference position, rank-out-of-bound flag, and token/rank lengths are useful localization evidence, but they cannot distinguish all interacting causes.

## 90 **Supplementary Note S5: Automated readability and model-based quality**

The computational stimulus design selected 504 rows balanced over seven conditions (72 per condition) and over all 18 prompt templates. Selection and randomization metadata were frozen in the retained manifest. Metrics were computed on visible

strings. Flesch Reading Ease is a syllable/sentence surface heuristic. Surface-flag counts search declared formatting and artifact patterns. TF-IDF cosine similarity measures lexical prompt overlap. None of these quantities records a reader’s judgment.



**Figure S2. Automated readability diagnostics.** Flesch Reading Ease and related surface diagnostics are summarized over 504 balanced computational stimuli, with uncertainty resampled over prompt templates. Segmented full messages include payload-bearing forced spans and non-payload natural tails. These measures are automated surface properties, not human ratings, evidence of human-perceived naturalness, or substitutes for a participant study; no participant was recruited and no rating was collected.

| Condition              | Mean Flesch | Prompt-template bootstrap 95% CI | Interpretation                      |
|------------------------|-------------|----------------------------------|-------------------------------------|
| Ordinary control       | 61.940597   | 54.897042–68.143167              | Automated control reference         |
| Segmented full message | 60.551834   | 56.873661–64.135677              | Forced spans plus non-payload tails |

**Table S9.** Principal automated readability rows. The complete figure includes all seven balanced conditions and associated surface diagnostics.

Source-model quality supplied 45 protocol/model contrasts; 42 had Holm-adjusted  $p < 0.05$ , with estimates from -5.298565 to +2.684536. The independently held-out evaluator also supplied 45 contrasts, 42 adjusted  $p < 0.05$ , from -6.161972 to +3.287497. Signs differed by protocol and model, so there is no universal quality ordering. The source-cover-log-probability and held-out-evaluator mixed models were singular. Their coefficient tables remain diagnostic; no fixed-effects fallback was substituted.

The empirical capacity–quality frontier shown in Main Fig. 2 contains 4,320 trials restricted to the four payload classes on the common hexadecimal support. Three models and six protocols form 18 model–protocol cells; separate forced-span and full-message views yield 36 plotted cells. For every cell, effective artifact bits per forced token use the same artifact information numerator, mean token log probability is reported on the vertical axis, and point area represents mean automated surface-flag burden. Intervals use 2,000 payload-bootstrap resamples. Non-hexadecimal payload classes are excluded by the common-support restriction, and no lead-in condition exists in the complete primary matrix.

Forced-span metrics average only payload-bearing token positions. Full-message metrics additionally include greedy tail positions and can therefore improve as the tail becomes longer even if the forced span is unchanged. Both are reported to avoid treating tail fluency as evidence about the embedded positions.

The human-study status manifest records recruitment authorized = false, survey deployed = false, participants = 0, and ratings = 0. Planning materials are retained only to document a possible future design. They were not administered and are not reproduced here as a completed instrument. A model-name mismatch in those planning materials (Llama 3.1 versus the executed Llama 3 model) further prevents treating them as an approved executed protocol.

## Supplementary Note S6: Ablations and topic effects

### Locked ablation extraction


The ablation matrix paired 48 payload groups within each available model. The canonical setting was no lead-in, eight ranks per segment, the safe-text filter, and dynamic semantic stopping. Differences in Table S10 are level minus canonical. These compact contrasts were extracted after outcomes were available, without fitting new models, and are classified as exploratory/post-outcome. Payload group is the pairing and 2,000-resample bootstrap unit. The printed table presents the reviewer-relevant contrasts; the complete 60-row machine-readable contrast table, including multiplicity fields and secondary outcomes, remains in the public repository.

Across the evaluated 2-, 4-, 8-, 16-, and 32-token conditions, lead-ins generally worsened likelihood, and longer lead-ins generally increased full-message length. The 32-token contrast had mean-log-probability difference -1.652327 (95% CI, -1.749091 to -1.553103) and full-token-count difference +159.590278 (125.829861 to 191.030035). This is an adverse boundary/context result, but the paired contrast is not a causal decomposition. Segment-size changes altered efficiency and cover length: moving from 8 to 32 ranks increased effective rate by 1.231461 bits per full token (1.073116 to 1.388188) and decreased full-message length by 174.500000 tokens (-208.405035 to -140.573264). Forced-token count is codec-determined and unchanged; the length difference arises from fewer segment tails.

Removing the filter changed mean log probability by -0.001797 (-0.037738 to 0.035394) and effective rate by -0.062969 (-0.134265 to 0.007651); both selected intervals include zero. The Mistral roundtrip-stable-filter condition was unavailable for 48 work units, so the stable-filter zero rows use two models, whereas the no-filter comparison uses three. These contrasts do not establish a filter benefit.

Relative to dynamic stopping, no tail increased effective rate by 3.171551 bits per full token and reduced full-message length by 254.951389 tokens; removing non-payload text cannot be interpreted as a cover-quality improvement. The fixed 20–60-token sentence-tail policy reduced full-message length by 65.708333 tokens. Dynamic stopping checks a frozen sentence-boundary/length predicate and then applies an emergency cap; it is not a reader judgment.

### Exploratory ablation contrasts



**Figure S3. Ablation summary.** Selected paired level-minus-canonical estimates are shown with payload-bootstrap 95% confidence intervals. Rows are exploratory/post-outcome evidence extraction from retained outputs. The Mistral roundtrip-stable-filter cell is unavailable for 48 work units; unavailable is not an execution failure. No-tail rate gains primarily reflect deletion of non-payload tokens and are not quality evidence.

| Factor and level       | $\Delta$ mean log probability (95% CI) | $\Delta$ full tokens (95% CI)             | $\Delta$ bits/full token (95% CI)   | Interpretation or availability                           |
|------------------------|----------------------------------------|-------------------------------------------|-------------------------------------|----------------------------------------------------------|
| Lead-in: 2 tokens      | -0.575880<br>[-0.667347, -0.491716]    | +13.534722<br>[-21.431424, 50.179861]     | -0.063224<br>[-0.157433, 0.024628]  | Likelihood lower; length and rate intervals include zero |
| Lead-in: 4 tokens      | -0.788830<br>[-0.883630, -0.694343]    | +0.520833<br>[-36.346181, 35.365104]      | -0.051850<br>[-0.130954, 0.029064]  | Likelihood lower; length and rate intervals include zero |
| Lead-in: 8 tokens      | -1.041886<br>[-1.115991, -0.961546]    | +56.750000<br>[25.909201, 87.697396]      | -0.179845<br>[-0.262452, -0.101276] | Lower likelihood, longer cover, lower rate               |
| Lead-in: 16 tokens     | -1.367867<br>[-1.450430, -1.285953]    | +124.645833<br>[85.268229, 164.115799]    | -0.299132<br>[-0.379449, -0.216430] | Lower likelihood, longer cover, lower rate               |
| Lead-in: 32 tokens     | -1.652327<br>[-1.749091, -1.553103]    | +159.590278<br>[125.829861, 191.030035]   | -0.388736<br>[-0.469659, -0.312450] | Largest adverse lead-in contrast                         |
| Segment size: 4 ranks  | +0.116065<br>[0.045376, 0.192279]      | +212.527778<br>[172.067535, 253.639236]   | -0.373965<br>[-0.451966, -0.300845] | More segment tails; longer cover and lower rate          |
| Segment size: 16 ranks | +0.064803<br>[-0.000267, 0.131660]     | -104.694444<br>[-135.813021, -74.589410]  | +0.524216<br>[0.418241, 0.640763]   | Fewer tails; shorter cover and higher rate               |
| Segment size: 32 ranks | +0.152306<br>[0.076589, 0.230426]      | -174.500000<br>[-208.405035, -140.573264] | +1.231461<br>[1.073116, 1.388188]   | Fewer tails; shortest selected cover contrast            |
| No filter              | -0.001797<br>[-0.037738, 0.035394]     | +30.090278<br>[1.998264, 59.432986]       | -0.062969<br>[-0.134265, 0.007651]  | No clear likelihood or rate benefit                      |
| Stable filter          | 0.000000<br>[0.000000, 0.000000]       | 0.000000<br>[0.000000, 0.000000]          | 0.000000<br>[0.000000, 0.000000]    | Two-model invariant; Mistral unavailable, not failed     |
| No tail                | 0.000000<br>[0.000000, 0.000000]       | -254.951389<br>[-292.821007, -219.249479] | +3.171551<br>[3.094005, 3.244380]   | Mechanical deletion of non-payload tokens                |
| Fixed sentence tail    | 0.000000<br>[0.000000, 0.000000]       | -65.708333<br>[-92.590451, -40.372917]    | -0.037921<br>[-0.113898, 0.033508]  | Shorter than dynamic; rate interval includes zero        |

**Table S10. Reviewer-relevant ablation contrasts.** Differences are level minus the canonical setting (no lead-in, eight ranks per segment, safe-text filter, and dynamic tail). Estimates pool three available models and 48 paired payloads per model except the stable-filter row, which contains two models because the 48-unit Mistral condition was unavailable. Intervals are payload-bootstrap 95% intervals. Mean log probability is evaluated on payload-bearing forced positions; zero tail-policy differences arise because tail choice does not change those positions. All rows are exploratory/post-outcome.


### Prompt-topic contrasts

The topic extraction compares the multi-topic and single-topic segmented schedules within model. Artifact counts are ratios on the response scale; the remaining estimates are multi-topic minus single-topic contrasts on their reported model scales. One of 15 rows excluded the null: the Llama expected artifact-count ratio was 1.276361 (95% CI, 1.114743–1.461409; Holm-adjusted  $p = 0.003294$ ). Every other Holm-adjusted value was at least 0.396248. This locked extraction is exploratory and post-outcome; it supports neither a general topic advantage nor a new confirmatory hypothesis.

A separate exploratory subgroup display defines tail gain as full-message mean token log probability minus forced-span mean token log probability. It contains 8,280 paired segments from 1,440 segmented trials across three models, six prompt categories, and single-topic and multi-topic schedules, giving 36 displayed groups. Every group contains 230 segments, and all 36 group means were positive, ranging from +2.825721 to +3.652576. The numbers of independent payload clusters differ by schedule: 40 for single-topic and 190 for multi-topic. Uncertainty uses 2,000-resample payload-cluster bootstrapping stratified by payload class. This subgroup estimand is not equivalent to the adjusted single-topic versus multi-topic mixed-model contrast; the adjusted source-quality contrasts ranged from -0.02214 to +0.01881 and had Holm-adjusted  $p = 1.0$  for all three models.

## Tail gain by topic

230 segments/group; payload clusters: single-topic n=40, multi-topic n=190



**Figure S4. Exploratory tail gain by prompt category.** Tail gain is full-message mean token log probability minus forced-span mean token log probability. The 8,280 paired segments arise from 1,440 segmented trials across three models, six prompt categories, and two topic schedules, yielding 36 displayed groups of 230 segments each. All group means are positive and range from +2.825721 to +3.652576. Uncertainty uses 2,000-resample payload-cluster bootstrapping stratified by payload class; independent payload-cluster counts are 40 for single-topic groups and 190 for multi-topic groups. This exploratory subgroup estimand is not the adjusted schedule contrast: adjusted source-quality contrasts range from -0.02214 to +0.01881, with Holm-adjusted  $p = 1.0$  for all three models.

**Table S11.** All 15 prompt-topic contrasts.

| Outcome                            | Model   | Estimate  | 95% CI                | $p_{\text{Holm}}$ |
|------------------------------------|---------|-----------|-----------------------|-------------------|
| Primary artifact count             | Llama   | 1.276361  | [1.114743, 1.461409]  | 0.00329358        |
| Primary artifact count             | Mistral | 1.015189  | [0.848773, 1.214233]  | 1                 |
| Primary artifact count             | Qwen    | 0.957147  | [0.781516, 1.172247]  | 1                 |
| Source-model log probability       | Llama   | -0.022140 | [-0.118718, 0.074439] | 1                 |
| Source-model log probability       | Mistral | 0.016347  | [-0.080231, 0.112925] | 1                 |
| Source-model log probability       | Qwen    | 0.018813  | [-0.077766, 0.115391] | 1                 |
| Effective artifact rate            | Llama   | 0.012473  | [-0.075124, 0.100070] | 1                 |
| Effective artifact rate            | Mistral | 0.006792  | [-0.080805, 0.094389] | 1                 |
| Effective artifact rate            | Qwen    | -0.015822 | [-0.103419, 0.071775] | 1                 |
| Held-out evaluator log probability | Llama   | -0.012801 | [-0.129427, 0.103825] | 1                 |
| Held-out evaluator log probability | Mistral | 0.012046  | [-0.104580, 0.128671] | 1                 |
| Held-out evaluator log probability | Qwen    | 0.024711  | [-0.091915, 0.141336] | 1                 |
| Payload throughput                 | Llama   | 0.018304  | [-0.016147, 0.052755] | 0.595439          |
| Payload throughput                 | Mistral | 0.026471  | [-0.007981, 0.060922] | 0.396248          |
| Payload throughput                 | Qwen    | -0.009378 | [-0.043829, 0.025073] | 0.595439          |

## Supplementary Note S7: Neural steganalysis and leakage sensitivity

### Corpus, architectures, and splits


The frozen corpus contains 15,840 rows: 7,920 RankCloak covers and 7,920 matched ordinary controls, grouped by 480 payloads. Payload groups, rather than rows, were assigned to train, validation, and test sets. The TextCNN-equivalent implementation lowercases and tokenizes text into a 30,000-term vocabulary, truncates or pads to 256 tokens, uses 300-dimensional learned embeddings and 100 filters at each of widths 3 and 5, applies dropout 0.5, and optimizes with Adam at learning rate 0.001. The pretrained DeBERTa-v3-base classifier uses maximum length 256, learning rate  $2 \times 10^{-5}$ , weight decay 0.01, and three epochs.

For each architecture, the frozen design comprises one matched split, 18 held-out-template splits, three leave-one-model splits, and six leave-one-codec splits. Thus 28 fits per architecture and 56/56 fits overall completed, producing 55,440 prediction rows. No row-level reassignment was used to improve performance.

Primary reported endpoints are ROC-AUC and balanced accuracy. Precision, Brier score, and true-positive rate at false-positive rate at most 1% are explicitly supplementary/exploratory post-freeze measures. All confidence intervals for an individual split use 2,000 payload-group bootstrap resamples. Ranges across heterogeneous held-out splits are descriptive minima and maxima, not confidence intervals. Every split-level estimate and interval for all retained metrics is supplied in machine-readable Supplementary Data Table S7.1 (`supplementary_tables/detector_split_metrics.csv`).

## Neural detectability

Higher = easier detection · D-F exploratory



**Figure S5. Full neural-detector evaluation.** The six panels report confirmatory frozen endpoints together with supplementary exploratory metrics. Matched error bars are payload-group bootstrap 95% confidence intervals. Held-out regime bars span the minimum and maximum across prespecified splits and are not confidence intervals. Brier score is lower-is-better; high values in the other panels indicate stronger distinguishability and therefore weaker concealment.

| Architecture       | Matched metric    | Estimate | Payload-bootstrap 95% CI | Status                    |
|--------------------|-------------------|----------|--------------------------|---------------------------|
| TextCNN-equivalent | ROC-AUC           | 0.990190 | 0.987740–0.992682        | Confirmatory frozen       |
| TextCNN-equivalent | Balanced accuracy | 0.951010 | 0.943456–0.958466        | Confirmatory frozen       |
| TextCNN-equivalent | Precision         | 0.954684 | 0.945891–0.963421        | Supplementary exploratory |
| TextCNN-equivalent | Brier score       | 0.039320 | 0.033497–0.045453        | Supplementary exploratory |
| TextCNN-equivalent | TPR at FPR ≤ 1%   | 0.840909 | 0.806517–0.889752        | Supplementary exploratory |
| DeBERTa-v3-base    | ROC-AUC           | 0.999834 | 0.999696–0.999939        | Confirmatory frozen       |
| DeBERTa-v3-base    | Balanced accuracy | 0.991919 | 0.987654–0.995665        | Confirmatory frozen       |
| DeBERTa-v3-base    | Precision         | 0.999487 | 0.998393–1.000000        | Supplementary exploratory |
| DeBERTa-v3-base    | Brier score       | 0.007895 | 0.004432–0.011787        | Supplementary exploratory |
| DeBERTa-v3-base    | TPR at FPR ≤ 1%   | 0.993434 | 0.988724–0.998387        | Supplementary exploratory |

**Table S12. Matched detector endpoints.** Higher ROC-AUC, balanced accuracy, precision, and low-FPR sensitivity are adverse for concealment; lower Brier score indicates better detector calibration.

| Architecture       | Regime (split count)   | ROC-AUC range     | Balanced-accuracy range |
|--------------------|------------------------|-------------------|-------------------------|
| TextCNN-equivalent | Held-out template (18) | 0.720421–0.981705 | 0.565909–0.919318       |
| TextCNN-equivalent | Leave one model (3)    | 0.925623–0.973650 | 0.850758–0.883333       |
| TextCNN-equivalent | Leave one codec (6)    | 0.966532–0.998634 | 0.895833–0.980556       |
| DeBERTa-v3-base    | Held-out template (18) | 0.996291–0.999979 | 0.952273–0.995455       |
| DeBERTa-v3-base    | Leave one model (3)    | 0.996772–0.999979 | 0.945455–0.970455       |
| DeBERTa-v3-base    | Leave one codec (6)    | 0.997917–1.000000 | 0.927778–0.998611       |

**Table S13.** Ranges across heterogeneous frozen held-out splits. These ranges are not pooled confidence intervals.

### Leakage, exclusion sensitivity, repeatability, and cost

Exact train/test visible-text identity and exact tokenizer-sequence identity checks passed in all 28 frozen split definitions. A separate normalized character  $n$ -gram diagnostic at its declared threshold found 53 train/test near-duplicate pairs, of which 44 crossed payload groups, affecting 16 splits. Excluding affected test payload groups without retraining changed ROC-AUC from -0.001925 to +0.004568 for TextCNN-equivalent and from -0.000038 to +0.000166 for DeBERTa-v3-base; other observed endpoint shifts were also small. This threshold-dependent sensitivity does not rebuild training data and therefore does not eliminate the design limitation.

Fixed-seed reruns on the same CUDA device produced exact metric equality for both architectures. This is evidence of same-device repeatability only. CPU neural training was infeasible in the retained run and no CPU result was substituted; the study makes no CPU/GPU equivalence claim. The representative CUDA benchmark recorded TextCNN wall/training times 91.358113/9.497936 seconds and peak allocated VRAM 306,247,680 bytes; DeBERTa values were 1,536.660910/1,425.063414 seconds and 6,380,851,200 bytes.

## Supplementary Note S8: Capacity–quality theory

The main Methods defines bounded forced length, nominal and effective rates, full-message overhead, and the same-context probability bound. This note gives finite-code consequences, proof conditions, and the retained-field audit without repeating those central equations.

### Finite code space and exact inversion

Let  $H$ ,  $B$ , and  $n_B$  have the meanings in Main Eq. (14). The unused capacity of the smallest fixed-length rank word is described by

$$\sigma_B = n_B \log_2 B - H, \quad \eta_B = \frac{2^H}{B^{n_B}} = 2^{-\sigma_B}, \quad u_B = 1 - \eta_B.$$

Here  $\sigma_B$  is code-space slack in bits,  $\eta_B$  is utilization, and  $u_B$  is the unused codeword fraction. For a power-of-two alphabet and an integer-length bitstream,  $\sigma_B$  equals literal terminal padding bits. The byte codec records those bits explicitly; for  $B = 8$ , one byte is the finite example  $n_B = 3$  with one terminal padding bit, while a multi-byte payload is packed as one continuous stream. Hex nibbles have zero slack when the represented source length is exactly four bits per canonical character.

The rank-bound proposition assumes that every forced history has at least  $B$  eligible candidates and that encoder and decoder reproduce the same model, token IDs, prompt context, mask, and ordering. A digit  $d \in \{0, \dots, B-1\}$  selects the token at filtered rank  $d+1$ ; exact replay ranks that same saved token as  $d+1$ . Serial induction over the forced span recovers the rank word, and injectivity of the codec recovers the representation. The proposition does not apply after an altered token, context, boundary, mask, or model, and it supplies no edit-correction property.

### Segmentation and overhead decomposition

For segment forced lengths  $s_j$ , lead-in lengths  $\ell_j$ , and tail lengths  $t_j$ , the saved records satisfy  $\sum_j s_j = n_B$ . The empirically audited residual can be decomposed as

$$D = N_{\text{full}} - n_B = \sum_{j=1}^J (\ell_j + t_j) + \delta, \quad \delta \geq 0.$$

Thus segmentation alone preserves forced-position count, while more resets can create more opportunities for non-payload continuations. A no-tail condition raises bits per full token by shrinking the denominator. This is an arithmetic consequence and not evidence that the resulting visible text is more readable or natural.

### Conditional likelihood gap

The per-history ordering in Main Eq. (16) permits a trace-level comparison with the rank-1 eligible token and, for bounded codecs, the rank- $B$  token. In addition to  $Q_B$ , the audit records

$$\Delta_B = Q_B - Q_{\text{greedy}} = \frac{1}{n_B} \sum_{i=1}^{n_B} [-\log p_M(y_i | h_i) + \log p_{(1)}(h_i)] \geq 0.$$

The upper rank- $B$  endpoint is evaluated only when its probability was retained under the identical filtered history. Direct subword traces have no single fixed  $B$ , so that endpoint is unavailable rather than imputed. These quantities are model-scale surprisals in nats per forced token, not linguistic or participant outcomes.


### Empirical checks and trace boundary

The validation inventory contained 17,424 eligible retained records. Of these, 7,008 had all fields required for complete capacity, rate, and supported quality-bound checks. Across evaluable traces there were zero forced-position-count violations, zero rank/rate-bound violations, and zero supported quality-bound violations. The nine aggregate cells in the theoretical-capacity panel of Supplementary Fig. S6 lay exactly on the theoretical-versus-observed forced-position identity. There were 2,976 positive cover-length residuals, attributable to natural tails or other counted cover extension.

The complete encoder–decoder cascade proposition is not directly trace-evaluable from the saved record structure: none of the retained rows simultaneously carries every intermediate object required to replay the entire proof chain as one trace. The algebraic results are therefore conditional proofs and the retained-field checks are supporting empirical validation, not a complete end-to-end trace proof. They also do not validate human naturalness.

This theoretical validation asks whether the retained fields satisfy the declared fixed-radix identities, rate bounds, and supported same-context likelihood bound. It is distinct from the empirical capacity–quality frontier in Main Fig. 2, which compares observed representation rates and automated model/surface diagnostics on common payload support.

Capacity and tail overhead



**Figure S6. Theoretical capacity-tail validation.** Fixed-radix payload length and effective full-message rate are checked against their declared algebraic bounds, while the cover-length residual isolates tokens beyond the forced payload positions. Across 17,424 retained records, 7,008 were evaluable for the complete plotted capacity and supported quality-bound checks; no forced-position, rate-bound, or evaluated quality-bound violation occurred. Positive residuals in 2,976 records represent natural tails or other cover overhead. This validation tests retained-field identities and differs from the empirical common-support frontier in Main Fig. 2; neither analysis is evidence of human-perceived naturalness.

| Validation item                    | Result | Interpretation                                      |
|------------------------------------|--------|-----------------------------------------------------|
| Eligible records                   | 17,424 | Broad retained validation inventory                 |
| Fully evaluable records            | 7,008  | All required trace fields present                   |
| Forced-position violations         | 0      | Observed count matched codec requirement            |
| Rank/rate violations               | 0      | Evaluated bounded-codec relations held              |
| Supported quality-bound violations | 0      | Model-scale trace bound held where evaluable        |
| Positive cover-length residuals    | 2,976  | Tail/cover overhead, not error                      |
| Complete cascade traces            | 0      | End-to-end proposition not directly trace-evaluable |


**Table S14.** Empirical capacity and traceability audit.

Supplementary Note S9: Computational overhead

Primary-stage timing used 12,144 matched payload rows across the 18 model–protocol cells in Table S15. Intervals are 2,000-resample payload-group bootstrap intervals. Generation time is retained inclusive wrapper wall time; representation, filter setup, model generation, and supported recovery fields are reported separately in the sealed tables but are not claimed to be perfectly isolated kernel timings. Encoding wrapper time and supported decoding measurements consequently should not be summed as independent hardware primitives.

Across the 18 cells, mean inclusive generation time ranged from 1.674368 to 16.087628 seconds and payload throughput from 13.631855 to 100.851772 bits/s; effective payload rate ranged from 0.710004 to 7.322003 bits per full token. Direct-rank

and hex-nibble protocols were generally faster than the byte-expanded conditions, while segmented protocols were generally slower because they repeated prompt and tail work. These retained inclusive wall-time measurements support computational accounting under the recorded hardware and harness, not universal hardware-performance claims.



**Figure S7. Complete computational-overhead summary.** Primary model–protocol cells show payload-bootstrap summaries for inclusive generation, encoding setup, supported recovery, throughput, and effective rate. The setup panel uses a logarithmic axis because cells span several orders of magnitude. CPU time and repeated warm-up measurements were unavailable; wall time is not relabeled as CPU time.

| Model   | Protocol               | <i>n</i> | Mean generation time, s (95% CI)   | Mean payload throughput, bit/s (95% CI) |
|---------|------------------------|----------|------------------------------------|-----------------------------------------|
| Llama   | Direct ranks           | 480      | 2.074349<br>(2.014282–2.135219)    | 87.658980<br>(85.892805–89.433262)      |
|         |                        |          | 6.776682<br>(6.572039–6.982913)    | 37.409987<br>(36.542507–38.322129)      |
|         | ASCII-16               | 480      | 8.896445<br>(8.619910–9.172562)    | 28.522326<br>(27.867278–29.213688)      |
|         | ASCII-8                | 480      | 3.236699<br>(3.099398–3.376963)    | 55.691381<br>(55.148301–56.225745)      |
|         | Hex-16                 | 240      | 15.946543<br>(14.850856–17.084494) | 13.827671<br>(13.155163–14.496873)      |
|         | Segmented multi-topic  | 240      | 16.087628<br>(15.017670–17.208756) | 13.631855<br>(12.954794–14.325096)      |
|         | Segmented single-topic | 240      |                                    |                                         |
| Qwen    | Direct ranks           | 480      | 2.632108<br>(2.565827–2.699101)    | 70.911677<br>(68.729360–73.171915)      |
|         |                        |          | 7.480319<br>(7.242905–7.714463)    | 33.843456<br>(33.153343–34.571271)      |
|         | ASCII-16               | 480      | 9.847052<br>(9.537974–10.155572)   | 25.724469<br>(25.198423–26.278772)      |
|         | ASCII-8                | 480      | 3.427555<br>(3.284745–3.575041)    | 52.584782<br>(52.074452–53.086870)      |
|         | Hex-16                 | 240      | 12.276814<br>(11.449256–13.133478) | 17.040206<br>(16.420206–17.655385)      |
|         | Segmented multi-topic  | 240      | 11.881962<br>(11.129651–12.646529) | 17.182540<br>(16.534412–17.783251)      |
|         | Segmented single-topic | 240      |                                    |                                         |
| Mistral | Direct ranks           | 480      | 1.674368<br>(1.637117–1.711610)    | 95.454287<br>(92.689931–98.318820)      |
|         |                        |          | 3.532546<br>(3.430360–3.637370)    | 71.264443<br>(69.655729–72.892828)      |
|         | ASCII-16               | 480      | 4.558496<br>(4.421249–4.699172)    | 55.340577<br>(54.149134–56.561483)      |
|         | ASCII-8                | 480      | 1.762009<br>(1.698523–1.826194)    | 100.851772<br>(99.168287–102.554774)    |
|         | Hex-16                 | 240      | 10.478402<br>(9.810233–11.196055)  | 20.368846<br>(19.484536–21.270824)      |
|         | Segmented multi-topic  | 240      | 10.858662<br>(10.074554–11.660671) | 20.014386<br>(19.095928–20.969663)      |
|         | Segmented single-topic | 240      |                                    |                                         |

**Table S15. Primary model–protocol inclusive generation time and payload throughput.** Llama denotes Meta-Llama-3-8B-Instruct, Qwen denotes Qwen2.5-7B-Instruct, and Mistral denotes Mistral-7B-Instruct-v0.3; all use the reported Q4\_K\_M configurations. Parentheses contain payload-bootstrap 95% confidence intervals.

| Detector           | Wall time (s) | Training time (s) | Peak allocated VRAM | Scope                    |
|--------------------|---------------|-------------------|---------------------|--------------------------|
| TextCNN-equivalent | 91.358113     | 9.497936          | 306,247,680 bytes   | One fixed CUDA benchmark |
| DeBERTa-v3-base    | 1,536.660910  | 1,425.063414      | 6,380,851,200 bytes | One fixed CUDA benchmark |

**Table S16.** Representative detector cost. CPU time, repeat warm-up measurements, and cross-device benchmarks were unavailable.

## Supplementary Note S10: Worked examples and claim-boundary audit

### Representative retained examples

**Payload to ranks to cover.** A retained Llama segmented-hex trial used the public SHA-256 test-vector serialization

7cec7f1e68bbef4fdf066a51ccbea4377ee369e1b53219743aa386615e12d44a.

The first eight encoded ranks were [8, 13, 15, 13, 8, 16, 2, 15]. Under the saved prompt/model/filter state, the associated forced token IDs were [1789, 198, 16, 627, 13617, 10652, 25, 9842], detokenizing to the deliberately mechanical span

For\nl.\nExample conversation: Write

The subsequent visible text

a friendly message to a friend about making relaxed plans for the weekend.

was a natural tail. It carried no payload rank and was ignored by the decoder. Saved-token replay reproduced all nibble ranks and the serialized digest exactly. This example illustrates deterministic rank realization, not an assertion that the forced span is human-natural.

**Visible-text retokenization failure.** In a retained ablation trial, detokenization followed by tokenizer reapplication first diverged in segment 6, forced position 5: expected token ID 11,623 at rank 8 became observed token ID 53,992 at replay rank 2. The visible string alone therefore did not preserve the token boundary needed by the saved-token contract.

**Transformation failure.** For the SHA-256 example above, straight-to-smart quote conversion failed at segment 1, position 1: expected token ID 1 at rank 9, while the transformed replay observed token ID 2,118 at rank 155. The first-divergence record localizes the mismatch but does not prove that a single changed glyph alone caused every later error.

## Claim-boundary audit

**Table S17.** Claim boundaries applied throughout the revision.

| Issue                | Supported statement                                                                                                                   | Prohibited or unresolved extension                                                                    |
|----------------------|---------------------------------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------|
| Exact-copy recovery  | 6,480/6,480 primary serialized payloads recovered from saved token IDs under matched replay                                           | Ordinary visible-text reliability; cross-model robustness                                             |
| Transmission         | Controlled frozen perturbations reveal severe fragility; tail-only truncation preserves recovery when only ignored tokens are removed | Broad edit robustness or general mitigation                                                           |
| Human evaluation     | 504 computational stimuli have automated surface/model diagnostics                                                                    | Human-perceived readability or naturalness; zero participants and ratings                             |
| Deployment           | No operational channel or deployment traffic was tested                                                                               | Practical covert communication or real-world validation                                               |
| Detectability        | Two neural architectures strongly distinguish evaluated covers; high AUC is adverse                                                   | Undetectability or resistance to future detectors                                                     |
| Multilingual scope   | 288/288 Spanish and 288/288 Simplified Chinese exact-copy secondary trials recovered                                                  | Broad multilingual naturalness or cross-language generalization                                       |
| Unavailable work     | 432 predeclared combinations lacked compatible retained outputs                                                                       | Calling unavailable combinations experimental failures                                                |
| Model replay         | Retained outputs are bound to model/configuration manifests                                                                           | Fresh local generation replay while configured GGUF files are absent                                  |
| Detector leakage     | Exact identity checks passed; near-duplicate sensitivity shifts were small                                                            | Elimination of the 53-pair/16-split lexical-overlap limitation                                        |
| Statistical models   | Reported conditional estimates retain diagnostic warnings                                                                             | Treating three singular mixed-model fits as unqualified evidence; substituting an undeclared fallback |
| Exploratory analyses | Compact ablation, topic, and auxiliary detector analyses are labeled post-outcome/exploratory where applicable                        | Recasting them as preregistered confirmatory evidence                                                 |
| Overhead             | Retained inclusive wall-time and fixed CUDA benchmark scopes are reported                                                             | CPU time, warm-up stability, CPU/GPU equivalence, or hardware-general performance                     |
| Theory               | Conditional rank/capacity relations and retained-field checks are supported                                                           | A directly trace-evaluated full cascade or human-quality proof                                        |
| Cryptography         | The carrier can encode outputs from declared cryptographic algorithms                                                                 | Cryptographic security supplied by RankCloak itself; confidentiality without an external cryptosystem |